

AGI is coming soon; the world isn't ready


The CEOs of Anthropic, OpenAI, and Google DeepMind; Nobel laureate Geoffrey Hinton; “godfather of AI” Yoshua Bengio; and hundreds of other AI scientists and executives have signed the following statement:

“Mitigating the risk of extinction from AI should be a global priority alongside other societal-scale risks such as pandemics and nuclear war.”

1. AI minds are *grown*, not engineered.

Traditional software


engineered, line by line



Behavior follows the specification.
Failures can be traced to specific lines.
Bugs are debugged and removed.
Objectives are explicit.

AI models

grown by optimization



Behavior emerges from training — no specification.
Failures cannot be traced; we cannot read the weights.
Bugs are retrained against — and may re-emerge.
Objectives are a proxy; the model pursues the training signal.

Every other piece of critical infrastructure has the column on the left. AI does not.

Modern AI is not engineered. No one wrote its behavior line by line; it's the result of optimization over trillions of parameters. The behavior *emerges*, and there is no specification, no traceable logic, no debugger.


Every other piece of safety-critical infrastructure (aircraft, medical devices, nuclear reactors) is built with the controls AI lacks.

The labs themselves cannot reliably predict or constrain what their own systems will do.



details →

2. Capability is doubling every ~3.5 months.



Source: METR Time Horizon 1.1 (May 2026). See details page.

METR measures the ability of models by comparing them to human completion times. **Since 2024 the time horizon has been doubling every ~3.5 months.** In just one year, AI ability will be >16× what it can do today.

Claude Mythos exists but is unreleased, because its cybersecurity capabilities are so powerful that Anthropic has restricted public access (Project Glasswing). The Opus 4.7 model now generally available was deliberately trained with reduced cyber capabilities.


These are not forecasts. They are measurements.



details →

see reverse for economic impacts, public opinion, and contact →

3. Massive CBRN and economic impacts: already starting.



Source: Kording & Marinescu (2025), Brookings. See details page.

American frontier labs evaluate every release on CBRN uplift: the degree to which the model could help a non-expert build chemical, biological, radiological, or nuclear weapons. The UK AI Security Institute reports that frontier models now outperform PhD-level experts on open-ended chemistry and biology questions.


Open-source models trail the frontier by ~6 months, with no safety guardrails. Whatever capability Anthropic, OpenAI, or DeepMind have today has a publicly available equivalent just 6 months later, available for anyone on Earth to use privately.

Programming is past the peak (75% usage already), and the economic impacts have begun.




details →

4. Voters want action: bipartisan, by wide margins.



Voters of every party think AI is moving too fast. 77% of Democrats, 69% of Independents, and 68% of Republicans agree.

80% are concerned about AI. 76% in House battleground districts want Congress-mandated pre-deployment safety testing. 3-to-1 oppose federal preemption of state AI laws, including a plurality of Trump voters.



Sources: Economist/YouGov (May 2026), FLI/YouGov (Oct 2025). See details page.

On superhuman AI specifically: 13-to-1 against developing as fast as possible. Among the most bipartisan policy preferences in any domain; neither party owns the issue yet.



details →

Daniel Parshall, Ph.D. · Canary Institute
dan@canaryinstitute.ai

Former physicist and data scientist. Coauthor, "Measuring AI's Economic Reach" (Canary Institute / GWU Center for Economic Research, 2026).

Canary Institute · AI policy research for the public interest



WEBSITE
canaryinstitute.ai



X
x.com/CanaryInst